

sheet4

November 9, 2020

```
[5]: import numpy as np
import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt
import sklearn as sk
from sklearn.model_selection import KFold
from sklearn.linear_model import Ridge, Lasso
```

```
[7]: data = pd.read_csv('Data_Q2', index_col=0)
data
```

```
[7]:
```

	x	y
0	1.624345	-1.209235
1	-0.611756	2.394638
2	-0.528172	2.528607
3	-1.072969	-5.347027
4	0.865408	4.262015
...
95	0.077340	4.646311
96	-0.343854	4.493099
97	0.043597	6.583315
98	-0.620001	1.881545
99	0.698032	7.117147

[100 rows x 2 columns]

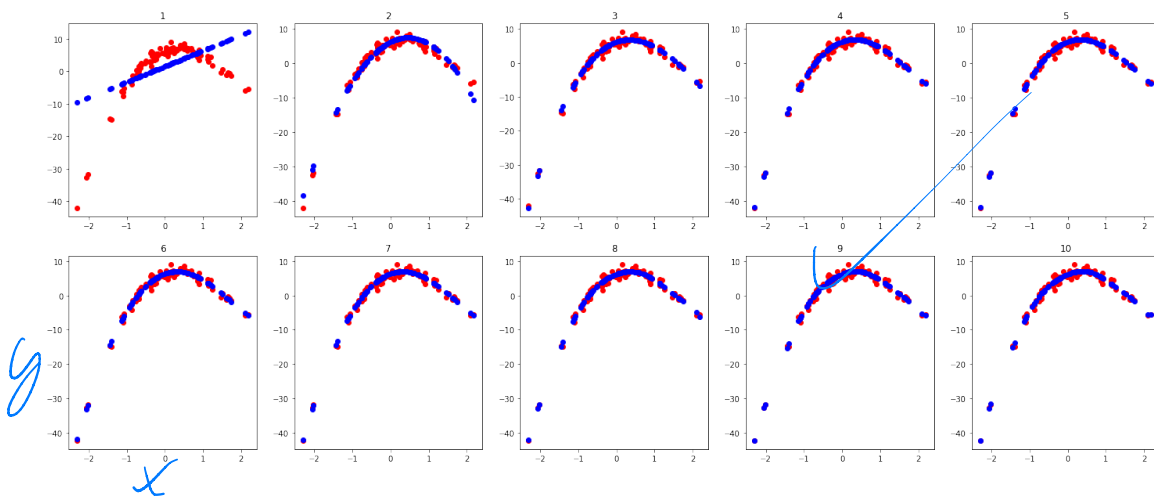
```
[9]: x = data['x']
y = data['y']
```

```
[10]: bic_vec = np.zeros((10,1))
aic_vec = np.zeros((10,1))
plt.figure(figsize=(25,10))
for i in range(10):
    p = i+1
    x_vec = np.polynomial.polynomial.polyvander(x,p)
    model = sm.OLS(y,x_vec).fit()
    if i==0:
```

```

        bic_min = model.bic
elif model.bic < bic_min:
    bic_min = model.bic
    best_bic_params = model.params
if i==0:
    aic_min = model.aic
elif model.aic < aic_min:
    aic_min = model.aic
    best_aic_params = model.params
bic_vec[i]=model.bic
aic_vec[i]=model.aic
plt.subplot(2,5,p)
y_predict = model.fittedvalues
plt.scatter(x,y,c='r')
plt.scatter(x,y_predict,c='b')
plt.title(str(p))

```



```

[11]: bic_aic_df = pd.DataFrame(np.concatenate((bic_vec,aic_vec),axis =1 ),columns =_
    ↳ ['bic value','aic value'], index = np.linspace(1,10,10))

```

```

[12]: bic_aic_df

```

```

[12]:      bic value  aic value
1.0    677.598182  672.387842
2.0    357.330237  349.514726
3.0    285.734557  275.313877
4.0    284.276388  271.250537
5.0    288.881554  273.250533
6.0    293.410946  275.174754
7.0    298.012627  277.171265

```

8.0	300.985570	277.539038
9.0	303.509673	277.457971
10.0	307.654182	278.997310

```
[13]: best_bic_params
```

```
[13]: const    6.313982
      x1      3.912735
      x2     -6.544477
      x3      0.699182
      x4      0.139361
      dtype: float64
```


```
[14]: best_aic_params
```

```
[14]: const    6.313982
      x1      3.912735
      x2     -6.544477
      x3      0.699182
      x4      0.139361
      dtype: float64
```

```
[15]: plt.figure()
      bic_aic_df.plot()
      plt.show()
```

<Figure size 432x288 with 0 Axes>

*Which one is best?
Write a short answer.*



0.1 part(e)

```
[16]: def ridge_lasso_reg2 (lambda_vec,option = 'ridge',params = False):

    kf = KFold(n_splits = 10)
    err = np.zeros(len(lambda_vec))
    X = np.polynomial.polynomial.polyvander(x,10)
    y = data['y']
    coefficients = np.zeros((len(lambda_vec),11))
    #     err1_vec = np.zeros(10)
    err2_vec = np.zeros(10)

    #carry out a binary search to find the optimal lambda
    for i in range(len(lambda_vec)):
        fold = 0
        if option == 'ridge':
            reg = Ridge(alpha = lambda_vec[i],fit_intercept = False)
        elif option == 'lasso':
            reg = Lasso(alpha = lambda_vec[i],fit_intercept = False,tol=0.
→01,max_iter = 10000)
        for train_index, test_index in kf.split(X):
            X_train, X_test = X[train_index],X[test_index]
            y_train, y_test = y[train_index],y[test_index]

        #         if i==0:
        #             ridreg1 = f(alpha = lambda_vec[0],fit_intercept = False,
→normalize=True)
        #             model1 = ridreg1.fit(X_train,y_train)
        #             y_pred_1 = model1.predict(X_test)
        #             err1_vec[fold] = np.mean((y_pred_1-y_test)**2)

        #     reg = f
        model = reg.fit(X_train,y_train)
        y_pred_2 = model.predict(X_test)
        err2_vec[fold] = np.mean((y_pred_2-y_test)**2)
        fold += 1

    #         if i==0:
    #             err1 = np.mean(err1_vec)
    #             err.append(err1)
    #         print(err2_vec)
    err2 = np.mean(err2_vec)
    err[i]=err2
    coefficients[i,:] = model.coef_
```

```

#         if i!=(n-1):
#             lambda_storage_vec.append(np.mean(lambda_vec))

#         if err1 > err2:
#             lambda_vec = [lambda_vec[1],np.mean(lambda_vec)]
#             err1 = err2
#         else:
#             lambda_vec = [lambda_vec[0],np.mean(lambda_vec)]

    if params:
        return err, coefficients
    else:
        return err

```

```

[17]: l1 = np.linspace(0,2,41)
ridge_err_2 = ridge_lasso_reg2(l1)
ridge_err_2

```

```

[17]: array([2.76115474, 3.05096039, 3.1041114 , 3.02655863, 2.88962884,
          2.73078843, 2.56947497, 2.41565579, 2.2741894 , 2.14711332,
          2.03489071, 1.93711259, 1.85290372, 1.78116113, 1.72069534,
          1.67031355, 1.6288677 , 1.59528087, 1.56856026, 1.54780208,
          1.53219122, 1.5209981 , 1.51357368, 1.50934374, 1.50780274,
          1.50850767, 1.51107219, 1.51516098, 1.52048452, 1.5267943 ,
          1.53387835, 1.54155725, 1.5496805 , 1.55812321, 1.56678314,
          1.57557807, 1.58444339, 1.59332995, 1.60220221, 1.61103646,
          1.61981935])

```

```


[18]: plt.plot(l1,ridge_err_2)

```

```

[18]: [<matplotlib.lines.Line2D at 0x2673baad848>]

```



we pick $\lambda = 1.2$

```
[22]: l2 = np.linspace(0,0.2,21) #no matter how I change this interval, it always
      → suggests that  $\lambda = 0$  gives the smallest value, I have also altered the tol
      → and max_iter a lot to make it converge, but no effect
lasso_err_2, lasso_params_2 = ridge_lasso_reg2(l2, option='lasso', params=True)
lasso_err_2
```

```
c:\users\xuais\anaconda3\envs\snakes\lib\site-packages\ipykernel_launcher.py:29:
UserWarning: With alpha=0, this algorithm does not converge well. You are
advised to use the LinearRegression estimator
c:\users\xuais\anaconda3\envs\snakes\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:531: UserWarning:
Coordinate descent with no regularization may lead to unexpected results and is
discouraged.
```

positive)

```
c:\users\xuais\anaconda3\envs\snakes\lib\site-packages\ipykernel_launcher.py:29:
UserWarning: With alpha=0, this algorithm does not converge well. You are
advised to use the LinearRegression estimator
c:\users\xuais\anaconda3\envs\snakes\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:531: UserWarning:
Coordinate descent with no regularization may lead to unexpected results and is
discouraged.
```

positive)

```
c:\users\xuais\anaconda3\envs\snakes\lib\site-packages\ipykernel_launcher.py:29:
UserWarning: With alpha=0, this algorithm does not converge well. You are
```

```

advised to use the LinearRegression estimator
c:\users\xuais\anaconda3\envs\snakes\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:531: UserWarning:
Coordinate descent with no regularization may lead to unexpected results and is
discouraged.
    positive)
c:\users\xuais\anaconda3\envs\snakes\lib\site-packages\ipykernel_launcher.py:29:
UserWarning: With alpha=0, this algorithm does not converge well. You are
advised to use the LinearRegression estimator
c:\users\xuais\anaconda3\envs\snakes\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:531: UserWarning:
Coordinate descent with no regularization may lead to unexpected results and is
discouraged.
    positive)
c:\users\xuais\anaconda3\envs\snakes\lib\site-packages\ipykernel_launcher.py:29:
UserWarning: With alpha=0, this algorithm does not converge well. You are
advised to use the LinearRegression estimator
c:\users\xuais\anaconda3\envs\snakes\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:531: UserWarning:
Coordinate descent with no regularization may lead to unexpected results and is
discouraged.
    positive)
c:\users\xuais\anaconda3\envs\snakes\lib\site-packages\ipykernel_launcher.py:29:
UserWarning: With alpha=0, this algorithm does not converge well. You are
advised to use the LinearRegression estimator
c:\users\xuais\anaconda3\envs\snakes\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:531: UserWarning:
Coordinate descent with no regularization may lead to unexpected results and is
discouraged.
    positive)
c:\users\xuais\anaconda3\envs\snakes\lib\site-packages\ipykernel_launcher.py:29:
UserWarning: With alpha=0, this algorithm does not converge well. You are
advised to use the LinearRegression estimator
c:\users\xuais\anaconda3\envs\snakes\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:531: UserWarning:
Coordinate descent with no regularization may lead to unexpected results and is
discouraged.
    positive)
c:\users\xuais\anaconda3\envs\snakes\lib\site-packages\ipykernel_launcher.py:29:
UserWarning: With alpha=0, this algorithm does not converge well. You are

```

```

advised to use the LinearRegression estimator
c:\users\xuais\anaconda3\envs\snakes\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:531: UserWarning:
Coordinate descent with no regularization may lead to unexpected results and is
discouraged.
    positive)
c:\users\xuais\anaconda3\envs\snakes\lib\site-packages\ipykernel_launcher.py:29:
UserWarning: With alpha=0, this algorithm does not converge well. You are
advised to use the LinearRegression estimator
c:\users\xuais\anaconda3\envs\snakes\lib\site-
packages\sklearn\linear_model\_coordinate_descent.py:531: UserWarning:
Coordinate descent with no regularization may lead to unexpected results and is
discouraged.
    positive)

```

```

[22]: array([ 0.93712915,  0.98974645,  1.87466238,  3.79763414,  4.31165054,
            4.72633177,  5.15139012,  5.59376147,  6.05532773,  6.48937656,
            6.99306678,  7.53412523,  8.14033929,  8.78599176,  9.47866918,
            10.23401447, 11.08399508, 12.06880469, 13.26868876, 14.27880193,
            15.19249394])

```

```

[23]: lasso_params_2

```

```

[23]: array([[ 6.02578958e+00,  3.99587237e+00, -5.94340822e+00,
            5.81034046e-01, -5.75622912e-02,  1.16659762e-02,
            8.37006692e-03,  1.30737347e-03,  1.22853896e-03,
            1.93059879e-04,  1.57053426e-04],
            [ 5.99479394e+00,  3.96747478e+00, -5.89318952e+00,
            5.85838642e-01, -6.68623683e-02,  1.21527270e-02,
            7.96396237e-03,  1.25206237e-03,  1.28194942e-03,
            1.90276926e-04,  1.64508548e-04],
            [ 5.96379830e+00,  3.93907719e+00, -5.84297081e+00,
            5.90643237e-01, -7.61624455e-02,  1.26394778e-02,
            7.55785782e-03,  1.19675127e-03,  1.33535988e-03,
            1.87493973e-04,  1.71963670e-04],
            [ 5.93280266e+00,  3.91067959e+00, -5.79275211e+00,
            5.95447833e-01, -8.54625226e-02,  1.31262287e-02,
            7.15175327e-03,  1.14144016e-03,  1.38877034e-03,
            1.84711020e-04,  1.79418792e-04],
            [ 5.92556842e+00,  3.88727699e+00, -5.77964731e+00,
            6.01223608e-01, -9.15185384e-02,  1.36135924e-02,
            7.13649811e-03,  1.05178623e-03,  1.50013443e-03,
            1.82532937e-04,  1.90009738e-04],
            [ 5.92662026e+00,  3.86444359e+00, -5.77747272e+00,
            6.08293061e-01, -9.77621508e-02,  1.42279069e-02,
            7.22112087e-03,  9.01969604e-04,  1.66985026e-03,
            1.80100490e-04,  2.03059105e-04],

```


[5.91304094e+00, 3.83739111e+00, -5.74952244e+00,
6.15476815e-01, -1.07835136e-01, 1.49997711e-02,
7.01738814e-03, 7.26569573e-04, 1.85257647e-03,
1.78037526e-04, 2.16829102e-04],

[5.89666494e+00, 3.80670088e+00, -5.71099809e+00,
6.24385351e-01, -1.22219866e-01, 1.62409373e-02,
6.63927566e-03, 4.16986955e-04, 2.14914702e-03,
1.77877331e-04, 2.33656076e-04],

[5.86413674e+00, 3.77334247e+00, -5.64173574e+00,
6.31817684e-01, -1.42100295e-01, 1.77821282e-02,
5.89583157e-03, 8.62691216e-05, 2.49249667e-03,
1.83730066e-04, 2.54079588e-04],

[5.82514889e+00, 3.74243652e+00, -5.56530793e+00,
6.36581416e-01, -1.60890757e-01, 1.88505224e-02,
5.14841808e-03, -0.00000000e+00, 2.74029142e-03,
1.75061355e-04, 2.74360341e-04],

[5.78321450e+00, 3.70942131e+00, -5.47616045e+00,
6.40836869e-01, -1.85158091e-01, 2.04962738e-02,
4.16591309e-03, -0.00000000e+00, 3.16074087e-03,
1.34985513e-04, 2.94303986e-04],

[5.73836944e+00, 3.67811602e+00, -5.37873913e+00,
6.43488606e-01, -2.12019763e-01, 2.23084953e-02,
3.10056097e-03, -0.00000000e+00, 3.62265310e-03,
1.00242917e-04, 3.18281037e-04],

[5.69137158e+00, 3.64815906e+00, -5.27465113e+00,
6.44739379e-01, -2.41294458e-01, 2.42792279e-02,
2.06849324e-03, -0.00000000e+00, 4.10684493e-03,
7.00922439e-05, 3.46390254e-04],

[5.63594809e+00, 3.62071922e+00, -5.14046681e+00,
6.41395015e-01, -2.81801346e-01, 2.74334729e-02,
4.44365173e-04, -0.00000000e+00, 4.91063477e-03,
3.37947113e-05, 3.81110817e-04],

[5.57602496e+00, 3.59700220e+00, -4.99397947e+00,
6.34023899e-01, -3.27738448e-01, 3.11883637e-02,
0.00000000e+00, 0.00000000e+00, 5.57984524e-03,
1.51687511e-08, 4.20736913e-04],

[5.53783225e+00, 3.56632811e+00, -4.92352277e+00,
6.39682194e-01, -3.44206961e-01, 3.14072096e-02,
0.00000000e+00, 0.00000000e+00, 5.57591155e-03,
-0.00000000e+00, 4.51873478e-04],

[5.49371500e+00, 3.53801132e+00, -4.83083865e+00,
6.40591591e-01, -3.69883308e-01, 3.26195187e-02,
0.00000000e+00, 0.00000000e+00, 5.75538811e-03,
-0.00000000e+00, 4.88360935e-04],

[5.44132390e+00, 3.51535417e+00, -4.70732108e+00,
6.35851281e-01, -4.08111375e-01, 3.49352588e-02,
0.00000000e+00, 0.00000000e+00, 6.16287903e-03,

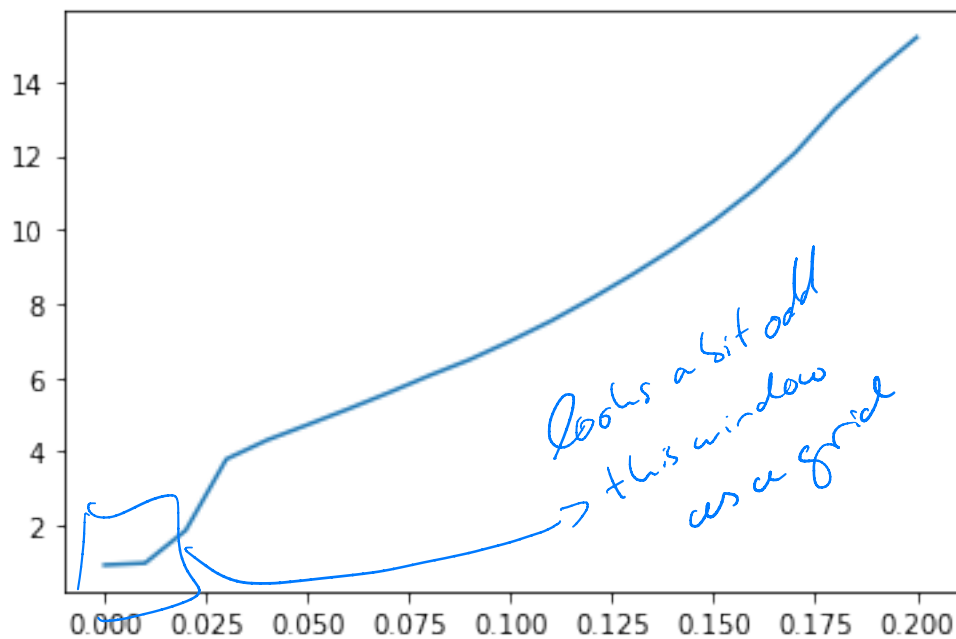
```

-0.00000000e+00,  5.35507539e-04],
[ 5.38769857e+00,  3.49288305e+00, -4.57955445e+00,
  6.30950301e-01, -4.48007774e-01,  3.72941334e-02,
  0.00000000e+00,  0.00000000e+00,  6.59457336e-03,
 -0.00000000e+00,  5.84996692e-04],
[ 5.33300588e+00,  3.47050018e+00, -4.44795755e+00,
  6.25818321e-01, -4.89440203e-01,  3.97175644e-02,
  0.00000000e+00,  0.00000000e+00,  7.05008433e-03,
 -0.00000000e+00,  6.36496175e-04],
[ 5.27738887e+00,  3.44838614e+00, -4.31298693e+00,
  6.20278047e-01, -5.32237320e-01,  4.22339945e-02,
  0.00000000e+00,  0.00000000e+00,  7.52847730e-03,
 -0.00000000e+00,  6.89506876e-04]]))

```

```
[24]: plt.plot(l2,lasso_err_2)
```

```
[24]: [<matplotlib.lines.Line2D at 0x2673ba8a508>]
```




```
[28]: X = np.polynomial.polynomial.polyvander(x,10)
```

```
[31]: reg = Lasso(alpha =0.5,tol=0.01)
model_lasso = reg.fit(X,y)
y_pred_lasso = model_lasso.predict(X)
plt.scatter(x,y,c='r')
plt.scatter(x,y_pred_lasso,c='b')
```


[31]: <matplotlib.collections.PathCollection at 0x2673b6c2f88>

Which model is that?



```
[32]: X = np.polynomial.polynomial.polyvander(x,10)
reg = Ridge(alpha =1.20)
model_ridge = reg.fit(X,y)
y_pred_ridge = model_ridge.predict(X)
plt.scatter(x,y,c='r')
plt.scatter(x,y_pred_ridge,c='b')
```

[32]: <matplotlib.collections.PathCollection at 0x2673b5d2d08>



[]: